

# Task VII: Equivariant Quantum Neural Networks

## Objective

This task involves implementing a  $Z_2 \times Z_2$  equivariant quantum neural network (EQNN) and comparing its performance against a **standard quantum neural network (QNN)**. The goal is to train these models on a dataset that respects the  $Z_2 \times Z_2$  symmetry, where data points are mirrored along  $y = x$ .

This implementation is based on the paper [arXiv:2205.06217](https://arxiv.org/abs/2205.06217) and additional background from [arXiv:2210.08566](https://arxiv.org/abs/2210.08566).

---

```
!pip install pennylane numpy torch matplotlib
```

```
import pennylane as qml
from pennylane import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
from torch.utils.data import Dataset, DataLoader, random_split
from sklearn.metrics import classification_report
from matplotlib.colors import ListedColormap
```

Collecting pennylane

```
  Downloading PennyLane-0.40.0-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: numpy in
/usr/local/lib/python3.11/dist-packages (2.0.2)
Requirement already satisfied: torch in
/usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: scipy in
/usr/local/lib/python3.11/dist-packages (from pennylane) (1.14.1)
Requirement already satisfied: networkx in
/usr/local/lib/python3.11/dist-packages (from pennylane) (3.4.2)
Collecting rustworkx<=0.14.0 (from pennylane)
  Downloading rustworkx-0.16.0-cp39-abi3-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (10 kB)
Requirement already satisfied: autograd in
/usr/local/lib/python3.11/dist-packages (from pennylane) (1.7.0)
Collecting tomlkit (from pennylane)
```

Downloading tomllkit-0.13.2-py3-none-any.whl.metadata (2.7 kB)  
Collecting appdirs (from pennylane)  
Downloading appdirs-1.4.4-py2.py3-none-any.whl.metadata (9.0 kB)  
Collecting autoray>=0.6.11 (from pennylane)  
Downloading autoray-0.7.1-py3-none-any.whl.metadata (5.8 kB)  
Requirement already satisfied: cachetools in  
/usr/local/lib/python3.11/dist-packages (from pennylane) (5.5.2)  
Collecting pennylane-lightning>=0.40 (from pennylane)  
Downloading PennyLane\_Lightning-0.40.0-cp311-cp311-  
manylinux\_2\_28\_x86\_64.whl.metadata (27 kB)  
Requirement already satisfied: requests in  
/usr/local/lib/python3.11/dist-packages (from pennylane) (2.32.3)  
Requirement already satisfied: typing-extensions in  
/usr/local/lib/python3.11/dist-packages (from pennylane) (4.12.2)  
Requirement already satisfied: packaging in  
/usr/local/lib/python3.11/dist-packages (from pennylane) (24.2)  
Collecting diastatic-malt (from pennylane)  
Downloading diastatic\_malt-2.15.2-py3-none-any.whl.metadata (2.6 kB)  
Requirement already satisfied: filelock in  
/usr/local/lib/python3.11/dist-packages (from torch) (3.18.0)  
Requirement already satisfied: jinja2 in  
/usr/local/lib/python3.11/dist-packages (from torch) (3.1.6)  
Requirement already satisfied: fsspec in  
/usr/local/lib/python3.11/dist-packages (from torch) (2025.3.0)  
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch)  
Downloading nvidia\_cuda\_nvrtc\_cu12-12.4.127-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch)  
Downloading nvidia\_cuda\_runtime\_cu12-12.4.127-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch)  
Downloading nvidia\_cuda\_cupti\_cu12-12.4.127-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch)  
Downloading nvidia\_cudnn\_cu12-9.1.0.70-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch)  
Downloading nvidia\_cublas\_cu12-12.4.5.8-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch)  
Downloading nvidia\_cufft\_cu12-11.2.1.3-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
Collecting nvidia-curand-cu12==10.3.5.147 (from torch)  
Downloading nvidia\_curand\_cu12-10.3.5.147-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.5 kB)  
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch)  
Downloading nvidia\_cusolver\_cu12-11.6.1.9-py3-none-  
manylinux2014\_x86\_64.whl.metadata (1.6 kB)  
Collecting nvidia-cuspars-cu12==12.3.1.170 (from torch)

```

    Downloading nvidia_cusparselt-cu12-12.3.1.170-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in
/usr/local/lib/python3.11/dist-packages (from torch) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in
/usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch)
    Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.2.0 in
/usr/local/lib/python3.11/dist-packages (from torch) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in
/usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch)
(1.3.0)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (4.56.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: pillow>=8 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (2.8.2)
Collecting scipy-openblas32>=0.3.26 (from pennylane-lightning>=0.40-
>pennylane)
    Downloading scipy_openblas32-0.3.29.0.0-py3-none-
manylinux_2_17_x86_64.whl.metadata (56 kB)
56.1/56.1 kB 4.6 MB/s eta
0:00:00
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
packages (from python-dateutil>=2.7->matplotlib) (1.17.0)
Requirement already satisfied: astunparse in
/usr/local/lib/python3.11/dist-packages (from diastatic-malt-
>pennylane) (1.6.3)
Requirement already satisfied: gast in /usr/local/lib/python3.11/dist-
packages (from diastatic-malt->pennylane) (0.6.0)
Requirement already satisfied: termcolor in
/usr/local/lib/python3.11/dist-packages (from diastatic-malt-
>pennylane) (2.5.0)
Requirement already satisfied: MarkupSafe>=2.0 in

```

```
/usr/local/lib/python3.11/dist-packages (from jinja2->torch) (3.0.2)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.11/dist-packages (from requests->pennylane)
(3.4.1)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.11/dist-packages (from requests->pennylane)
(3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.11/dist-packages (from requests->pennylane)
(2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.11/dist-packages (from requests->pennylane)
(2025.1.31)
Requirement already satisfied: wheel<1.0,>=0.23.0 in
/usr/local/lib/python3.11/dist-packages (from astunparse->diastatic-
malt->pennylane) (0.45.1)
Downloading PennyLane-0.40.0-py3-none-any.whl (2.0 MB)
_____ 2.0/2.0 MB 32.7 MB/s eta
0:00:00
anylinux2014_x86_64.whl (363.4 MB)
_____ 363.4/363.4 MB 3.8 MB/s eta
0:00:00
anylinux2014_x86_64.whl (13.8 MB)
_____ 13.8/13.8 MB 58.3 MB/s eta
0:00:00
anylinux2014_x86_64.whl (24.6 MB)
_____ 24.6/24.6 MB 43.2 MB/s eta
0:00:00
e_cul2-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
_____ 883.7/883.7 kB 36.4 MB/s eta
0:00:00
anylinux2014_x86_64.whl (664.8 MB)
_____ 664.8/664.8 MB 2.5 MB/s eta
0:00:00
anylinux2014_x86_64.whl (211.5 MB)
_____ 211.5/211.5 MB 4.9 MB/s eta
0:00:00
anylinux2014_x86_64.whl (56.3 MB)
_____ 56.3/56.3 MB 12.5 MB/s eta
0:00:00
anylinux2014_x86_64.whl (127.9 MB)
_____ 127.9/127.9 MB 6.9 MB/s eta
0:00:00
anylinux2014_x86_64.whl (207.5 MB)
_____ 207.5/207.5 MB 5.2 MB/s eta
0:00:00
anylinux2014_x86_64.whl (21.1 MB)
_____ 21.1/21.1 MB 66.1 MB/s eta
0:00:00
```

```

930.8/930.8 kB 47.5 MB/s eta
0:00:00
anylinux_2_28_x86_64.whl (2.4 MB)
2.4/2.4 MB 66.9 MB/s eta
0:00:00
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.1 MB)
2.1/2.1 MB 56.6 MB/s eta
0:00:00
alt-2.15.2-py3-none-any.whl (167 kB)
167.9/167.9 kB 13.0 MB/s eta
0:00:00
lkit-0.13.2-py3-none-any.whl (37 kB)
Downloading scipy_openblas32-0.3.29.0.0-py3-none-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (8.6 MB)
8.6/8.6 MB 82.1 MB/s eta
0:00:00
lkit, scipy-openblas32, rustworkx, nvidia-nvjitlink-cu12, nvidia-
curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-
nVRTC-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, autoray,
nvidia-cusparse-cu12, nvidia-cudnn-cu12, diastatic-malt, nvidia-
cusolver-cu12, pennylane-lightning, pennylane
Attempting uninstall: nvidia-nvjitlink-cu12
Found existing installation: nvidia-nvjitlink-cu12 12.5.82
Uninstalling nvidia-nvjitlink-cu12-12.5.82:
Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
Attempting uninstall: nvidia-curand-cu12
Found existing installation: nvidia-curand-cu12 10.3.6.82
Uninstalling nvidia-curand-cu12-10.3.6.82:
Successfully uninstalled nvidia-curand-cu12-10.3.6.82
Attempting uninstall: nvidia-cufft-cu12
Found existing installation: nvidia-cufft-cu12 11.2.3.61
Uninstalling nvidia-cufft-cu12-11.2.3.61:
Successfully uninstalled nvidia-cufft-cu12-11.2.3.61
Attempting uninstall: nvidia-cuda-runtime-cu12
Found existing installation: nvidia-cuda-runtime-cu12 12.5.82
Uninstalling nvidia-cuda-runtime-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cu12
Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cu12
Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
Found existing installation: nvidia-cublas-cu12 12.5.3.2
Uninstalling nvidia-cublas-cu12-12.5.3.2:
Successfully uninstalled nvidia-cublas-cu12-12.5.3.2

```

```
Attempting uninstall: nvidia-cusparse-cu12
Found existing installation: nvidia-cusparse-cu12 12.5.1.3
Uninstalling nvidia-cusparse-cu12-12.5.1.3:
  Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3
Attempting uninstall: nvidia-cudnn-cu12
Found existing installation: nvidia-cudnn-cu12 9.3.0.75
Uninstalling nvidia-cudnn-cu12-9.3.0.75:
  Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
  Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed appdirs-1.4.4 autoray-0.7.1 diastatic-malt-
2.15.2 nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127
nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127
nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-curand-
cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-cu12-
12.3.1.170 nvidia-nvjitlink-cu12-12.4.127 pennylane-0.40.0 pennylane-
lightning-0.40.0 rustworkx-0.16.0 scipy-openblas32-0.3.29.0.0 tomlkit-
0.13.2
```

## Dataset Generation

A synthetic dataset was generated with **two features** ( $(x_1, x_2)$ ) and **two classes (0 and 1)**. The class label is determined based on the quadrant of the point:

- $(x_1, x_2)$  belongs to **Class 1** if  $(x_1 \times x_2 > 0)$
- Otherwise, it belongs to **Class 0**

To enforce  **$Z_2 \times Z_2$  symmetry**, transformations were applied:

- Mirroring along  $y = x$
- Mirroring along  $x = -x$
- Mirroring along  $y = -y$

This created a larger dataset where the labels remained consistent under these transformations.

---

```
# Set random seed for reproducibility
np.random.seed(42)

# Generate  $Z_2 \times Z_2$  symmetric dataset
def generate_symmetric_data(n_samples=100):
    """Generate a classification dataset that is  $Z_2 \times Z_2$  symmetric."""
    X = np.random.uniform(-1, 1, (n_samples, 2)) # Random points (x1,
x2)
    Y = np.array([1 if x[0] * x[1] > 0 else 0 for x in X]) # Label
```

*based on quadrant*

```
# Apply symmetry transformations (mirroring along y=x)
X_sym = np.vstack([X, X[:, ::-1], -X, -X[:, ::-1]]) # Generate
all symmetric variants
Y_sym = np.tile(Y, 4) # Labels remain the same

return X_sym, Y_sym

# Custom Dataset Class
class SymmetricDataset(Dataset):
    def __init__(self, X, Y):
        self.X = torch.tensor(X, dtype=torch.float32)
        self.Y = torch.tensor(Y, dtype=torch.float32)

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return self.X[idx], self.Y[idx]

# Prepare dataset
X, Y = generate_symmetric_data(10000)
dataset = SymmetricDataset(X, Y)

# Split dataset into train and test sets
train_size = int(0.8 * len(dataset))
test_size = len(dataset) - train_size
train_dataset, test_dataset = random_split(dataset, [train_size,
test_size])

# Create DataLoaders
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

## Quantum Neural Networks (QNN)

### Standard QNN

- Used **Angle Embedding** to encode input features
- Applied **Strongly Entangling Layers** for variational learning
- The output was measured using **PauliZ expectation values**

### Equivariant QNN (EQNN)

- Maintains  $Z_2 \times Z_2$  symmetry throughout the circuit

- Uses **parameterized rotation gates** (RY, RZ)
  - Applies **CNOT gates** to enforce equivariance
- 

```
# Quantum device setup
n_qubits = 2
dev = qml.device("default.qubit", wires=n_qubits)

# Standard QNN Circuit
def standard_qnn(weights, inputs):
    """Standard variational quantum circuit."""
    qml.AngleEmbedding(inputs, wires=[0, 1])
    qml.StronglyEntanglingLayers(weights, wires=[0, 1])
    return qml.expval(qml.PauliZ(0))

# QNode Wrapper
qnode_standard = qml.QNode(standard_qnn, dev, interface="torch")

# PyTorch Model Wrapper
class QuantumModel(nn.Module):
    def __init__(self, qnode, n_layers=2):
        super().__init__()
        weight_shapes = {"weights": (n_layers, n_qubits, 3)}
        self.q_layer = qml.qnn.TorchLayer(qnode, weight_shapes)

    def forward(self, x):
        return self.q_layer(x)

# Instantiate Standard Model
standard_model = QuantumModel(qnode_standard, n_layers=5)

# Initialize weights properly
for param in standard_model.parameters():
    nn.init.uniform_(param, -np.pi, np.pi)

# Loss Function and Optimizer
criterion = nn.BCEWithLogitsLoss()
optimizer_std = optim.Adam(standard_model.parameters(), lr=0.001)
scheduler_std = optim.lr_scheduler.StepLR(optimizer_std, step_size=10,
gamma=0.5)

# Training Function
def train_model(model, optimizer, scheduler, train_loader,
epochs=500):
    model.train()
    losses = []
    for epoch in range(epochs):
        epoch_loss = 0
```



```

correct, total = 0, 0
for X_batch, Y_batch in train_loader:
    optimizer.zero_grad()
    outputs = model(X_batch).squeeze()
    loss = criterion(outputs, Y_batch)
    loss.backward()
    optimizer.step()
    epoch_loss += loss.item()

    # Compute accuracy
    predictions = (torch.sigmoid(outputs) > 0.5).float()
    correct += (predictions == Y_batch).sum().item()
    total += Y_batch.size(0)

scheduler.step()
accuracy = correct / total
losses.append(epoch_loss / len(train_loader))

if epoch % 5 == 0:
    print(f"Epoch {epoch}, Loss: {losses[-1]:.4f}, Accuracy:
{accuracy:.4f}")

return losses

# Evaluation Function with Metrics
def evaluate_model_metrics(model, test_loader):
    model.eval()
    all_preds, all_labels = [], []
    with torch.no_grad():
        for X_batch, Y_batch in test_loader:
            outputs = model(X_batch).squeeze()
            predictions = (torch.sigmoid(outputs) > 0.5).float()
            all_preds.extend(predictions.cpu().numpy())
            all_labels.extend(Y_batch.cpu().numpy())

    print(classification_report(all_labels, all_preds, digits=4))

# Train & Evaluate Standard Model
losses = train_model(standard_model, optimizer_std, scheduler_std,
train_loader, epochs=25)
evaluate_model_metrics(standard_model, test_loader)

```

```

Epoch 0, Loss: 0.6440, Accuracy: 0.7621
Epoch 5, Loss: 0.5976, Accuracy: 0.9571
Epoch 10, Loss: 0.5975, Accuracy: 0.9664
Epoch 15, Loss: 0.5975, Accuracy: 0.9599
Epoch 20, Loss: 0.5975, Accuracy: 0.9726

```

	precision	recall	f1-score	support
0.0	0.9561	1.0000	0.9776	3945

	1.0	1.0000	0.9554	0.9772	4055
accuracy				0.9774	8000
macro avg		0.9781	0.9777	0.9774	8000
weighted avg		0.9784	0.9774	0.9774	8000

*# Equivariant QNN Circuit*

```
def equivariant_qnn(weights, inputs):
    """Equivariant variational quantum circuit."""
    qml.AngleEmbedding(inputs, wires=[0, 1])
```

```
    for i in range(n_qubits):
        qml.RY(weights[i], wires=i)
        qml.RZ(weights[i + n_qubits], wires=i)
```

```
    qml.CNOT(wires=[0, 1])
    qml.CNOT(wires=[1, 0])
```

```
    for i in range(n_qubits):
        qml.RY(weights[i + 2 * n_qubits], wires=i)
        qml.RZ(weights[i + 3 * n_qubits], wires=i)
```

```
    return qml.expval(qml.PauliZ(0))
```

*# QNode for Equivariant QNN*

```
dev_eq = qml.device("default.qubit", wires=n_qubits)
qnode_equivariant = qml.QNode(equivariant_qnn, dev_eq,
interface="torch")
```

*# PyTorch Model Wrapper for EQNN*

```
class EquivariantQuantumModel(nn.Module):
    def __init__(self, qnode):
        super().__init__()
        weight_shapes = {"weights": (4 * n_qubits,)}
        self.q_layer = qml.qnn.TorchLayer(qnode, weight_shapes)
```

```
    def forward(self, x):
        return self.q_layer(x)
```

*# Instantiate Equivariant Model*

```
equivariant_model = EquivariantQuantumModel(qnode_equivariant)
```

*# Initialize weights properly*

```
for param in equivariant_model.parameters():
    nn.init.uniform_(param, -np.pi, np.pi)
```

*# Optimizer and Scheduler for EQNN*

```
optimizer_eq = optim.Adam(equivariant_model.parameters(), lr=0.001)
scheduler_eq = optim.lr_scheduler.StepLR(optimizer_eq, step_size=10,
```

```
gamma=0.5)
```

```
# Train & Evaluate Equivariant QNN
```

```
losses_eq = train_model(equivariant_model, optimizer_eq, scheduler_eq,  
train_loader, epochs=25)
```

```
evaluate_model_metrics(equivariant_model, test_loader)
```

```
Epoch 0, Loss: 0.6210, Accuracy: 0.7920
```

```
Epoch 5, Loss: 0.5975, Accuracy: 0.9828
```

```
Epoch 10, Loss: 0.5974, Accuracy: 0.9832
```

```
Epoch 15, Loss: 0.5974, Accuracy: 0.9872
```

```
Epoch 20, Loss: 0.5974, Accuracy: 0.9938
```

	precision	recall	f1-score	support
0.0	0.9719	1.0000	0.9858	3945
1.0	1.0000	0.9719	0.9857	4055
accuracy			0.9858	8000
macro avg	0.9860	0.9859	0.9857	8000
weighted avg	0.9862	0.9858	0.9857	8000

```
# Plot Loss Curve Comparison
```

```
plt.plot(losses_eq, label="Equivariant QNN")
```

```
plt.plot(losses, label="Standard QNN")
```


```
plt.xlabel("Epochs")
```

```
plt.ylabel("Loss")
```

```
plt.title("Training Loss Comparison")
```

```
plt.legend()
```

```
plt.show()
```




```
# Decision Boundary Visualization
def plot_decision_boundary(model, title="Decision Boundary"):
    x_min, x_max = -1.2, 1.2
    y_min, y_max = -1.2, 1.2
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100),
                          np.linspace(y_min, y_max, 100))
    grid = torch.tensor(np.c_[xx.ravel(), yy.ravel()],
                        dtype=torch.float32)

    with torch.no_grad():
        preds = model(grid).squeeze()
        preds = (torch.sigmoid(preds) > 0.5).float().numpy()

    plt.contourf(xx, yy, preds.reshape(xx.shape), alpha=0.6)
    plt.scatter(X[:, 0], X[:, 1], c=Y, edgecolor='k')
    plt.title(title)
    plt.show()

plot_decision_boundary(equivariant_model, "Equivariant QNN Decision Boundary")
```



## Training and Evaluation

Both models were trained using **Binary Cross-Entropy Loss** with the **Adam optimizer** over **25 epochs**. The dataset was split into **80% training** and **20% testing**, using **batch size = 32**.

---

## Results

Model	Accuracy (Train)	Accuracy (Test)	F1-Score
Standard QNN	97.26%	97.74%	0.9774
Equivariant QNN	99.38%	98.58%	0.9857

## Key Observations

- The **Equivariant QNN** achieved **higher accuracy** and **faster convergence** compared to the **Standard QNN**.
  - The EQNN reached **98.28% accuracy in just 5 epochs**, compared to the Standard QNN which took longer.
  - The **F1-score** of the **Equivariant QNN** (0.9857) is slightly better than that of the **Standard QNN** (0.9774), indicating improved classification performance.
-

## Conclusion

This experiment demonstrates the advantage of incorporating symmetry constraints in quantum machine learning models. The **Equivariant QNN** outperforms the **Standard QNN** in both accuracy and convergence speed, proving the effectiveness of symmetry-aware architectures in quantum neural networks.

